



# 기계이상진단을 위한 인공지능 학습 기법

## 제 6강 데이터 증강과 이상치 노출

한국과학기술원 전기및전자공학부

최정우

jwoo@kaist.ac.kr

**KAIST EE**

# 목차

- 데이터 전처리 기법
  - 데이터 증강 기법 (data augmentation)
- 이상진단을 위한 분류 태스크 성능 향상 기법
  - Outlier Exposure
  - Mixup


# Lack of training data

- New production model
  - Human resources are required for data collection
  - High-variability of normal data depending on the operating condition
- Insufficiency in machine condition measurements
  - One cannot measure enough data before the activation of a new plant




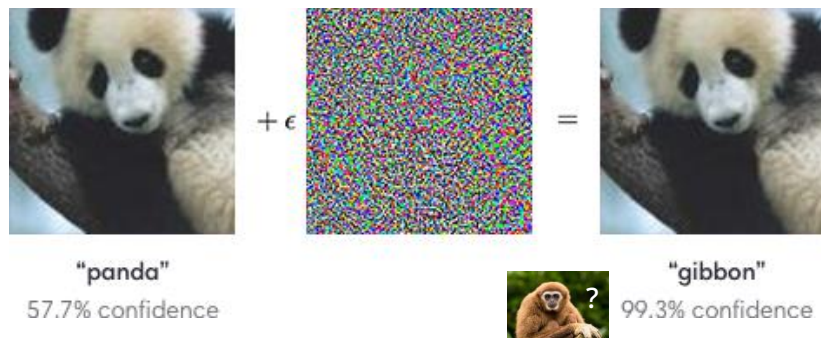
# Lack of training data

- In self-supervised training



# Example of overfitted model


- Adversarial attack examples
  - Adding small noise



Images from  
<https://arxiv.org/pdf/1602.02697.pdf>

# MNIST example


- Training dataset: 600 samples (out of 24,000 samples)
  - Normal: digits 0 – 4
  - Anomalies: digits 7,8,9




- Model
  - Pretext: classification task
  - Conv 2D classification model (3 conv 2D layers + one linear)
  - Anomaly score: maximum softmax probability
  - Label smoothing: 0.2
  - Test over 3584 samples

# Problems with small dataset

- Overfitting





- Big gap between performances for the training & validation datasets
- Model only works well for the training data




# Problems with small dataset

- Low AD performance




Even cannot recognize  
(unseen) normal data


original  
8




original  
7




original  
7




original  
9




original  
0




original  
7




original  
3




original  
0




original  
8



original  
8




original  
1






# Problems with small dataset

- Scattered samples in feature space




# Augmentation

- Increasing the size of training data
  - Using proper transforms




- Image: translation, rotation, flip, zoom, ...

- What is a 'proper transform'?
  - Relies on the **human's prior knowledge**
  - $T(\text{normal data}) \in \text{normal data}$



# Data Augmentation

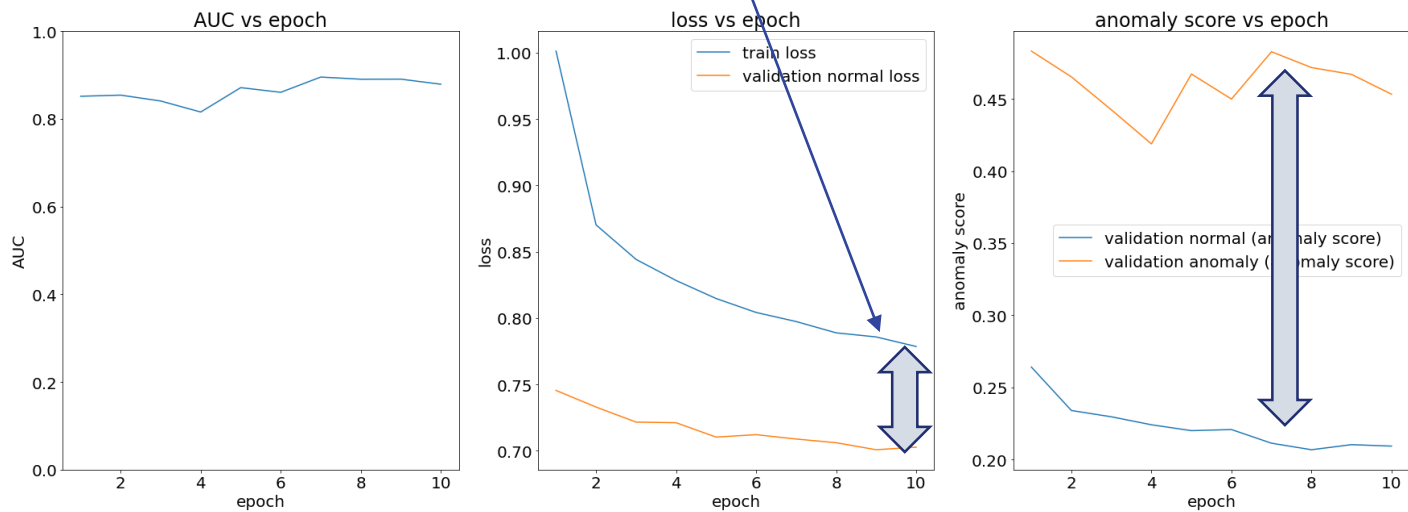
- Examples of augmented data (rotation  $< \pm 30^\circ$ , pixel shift  $< \pm 2$  px)



# Training with augmented dataset

- Training performance



smaller gap between train & validation losses



- Smaller gap between training & validation losses
- Anomaly scores for normal & abnormal data are more distinguishable

# Training with augmented dataset

- High AD performance



original  
8

original  
1

original  
8

original  
1

original  
9

original  
0


original  
3

original  
0

original  
9


original  
0

...



# Training with augmented dataset

- More separation in feature space



# Code review

- Augmentation part



```
# define transforms for augmentation
```

```
tfs = TF.Compose([  
    TF.RandomAffine(degrees=0, translate=(0.1,0.1)),  
    TF.RandomRotation(degrees=(-30, 30))  
])
```

← declare augmentation transform object

```
# get small sized dataset
```

```
# Sample the data in same proportion of labels
```

```
from sklearn.model_selection import StratifiedShuffleSplit
```

```
sss = StratifiedShuffleSplit(n_splits=1, train_size=train_dataset_num, random_state=0)
```

← Sample 600 data from MNIST train dataset

```
indices = list(range(len(mnist_train))) # gen list 0, ... ,len(mnist_train)
```

```
labels = [y for _, y in mnist_train] # get labels from mnist
```

```
train_idx = [idx for idx, _ in sss.split(indices, labels)] # separate indices
```

```
train_subset = Subset(mnist_train, train_idx[0]) # get subset
```

# Code review

- Augmentation part

```
# Begin augmentation
# init dataset list with original dataset
aug_dataset = [getSubset(train_subset)]

if not NAUG==0:
    # repeat random augmentation NAUG-1 times
    for ii in range(NAUG-1):
        aug_dataset.append(getSubset(train_subset, transform=tfs )) ← Apply augmentation and attach to list

# concatenate datasets using ConcatDataset function
train_dataset = torch.utils.data.ConcatDataset(aug_dataset) ← Unify datasets using ConcatDataset

train_dataloader = DataLoader(train_dataset, batch_size=BATCH, shuffle=True)
```



# Code review

- getSubset: customized class for building sub-dataset with transforms

```
class getSubset(TensorDataset):  
    def __init__(self, subset, transform=None, target_transform=None):  
        self.subset = subset  
        self.transform = transform  
        self.target_transform = target_transform  
  
    def __getitem__(self, index):  
        x, y = self.subset[index]  
        if self.transform:  
            x = self.transform(x)  
        if self.target_transform:  
            y = self.target_transform(y)  
        return x, y  
  
    def __len__(self):  
        return len(self.subset)
```

← get the transform for data

← get the transform for labels


← extract subset

← Apply transform to data


← Apply transform to labels

# Limitation of Data Augmentation

- Augmentation is not a fundamental solution (rather a sidekick)
  - Transformed normal data are still normal
  - Cannot make complex modifications: possible expressions are limited



- Some transforms can yield confusion between labels




# Outlier exposure

Injection of outlier examples to model

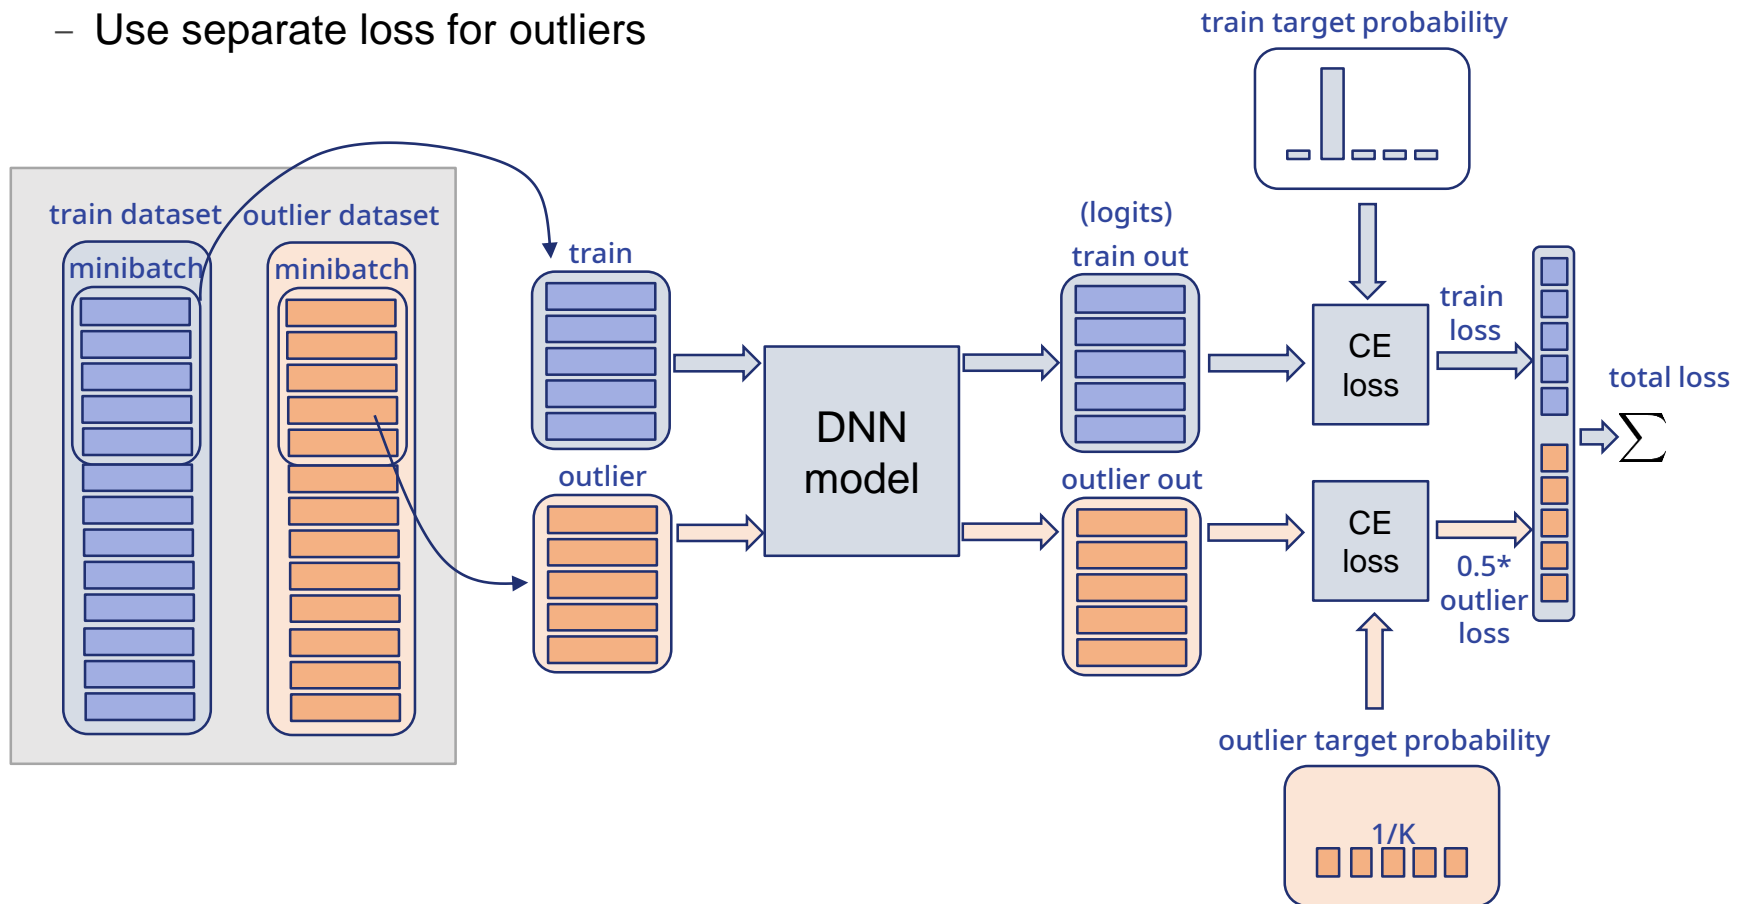
# Outlier exposure

- Prior knowledge of anomalies
  - Some data **can never be normal**
  - Exposing a model to **out-of-distribution (OOD) examples** entirely disjoint from normal data (auxiliary data)
  - Learning **cues** for whether an input is **modeled / unmodeled**
- Benefits
  - Learn effective heuristics for detecting OOD samples
  - Enabling the detection of novel forms of anomalies



# Outlier exposure

- How to
  - Train normal & outlier data together in the same batch
  - Use separate loss for outliers



# Outlier exposure

- Loss function

$$\mathbb{E}_{(x) \sim \mathcal{D}_{\text{in}}} [\mathcal{L}(f(x))] + \lambda \mathbb{E}_{x' \sim \mathcal{D}_{\text{out}}^{\text{OE}}} [\mathcal{L}_{\text{OE}}(f(x'))]$$


in-distribution  
(normal data)

out-of-distribution (OOD)  
(outliers)

- $\lambda$ : controls relative weight between two losses

- Example: MNIST

- Convolutional Classifier: same model
- Label smoothing: 0.2
- No augmentation, full training set



# Outlier exposure: code review

- Hyperparameters

```
▶ EPOCHS = 5           # Number of epochs to train
  BATCH = 32           # Minibatch size
  ORGCLASS_NUM = 10    # Num of original classes (10: 0 to 9)
  ANOMALY_NUM = [7,8,9] # (list) Digits will be used as anomalous data
  NORMAL_NUM = [0,1,2,3] # (list) Digits used as normal data
  OUTLIER_NUM = [4,5,6] # (list) Digits used as outlier

  TEMPSC = 1.0         # temperature parameter (for temperature scaling)
  LBSMOOTH = 0.2        # label smoothing parameter

  SCORE_MODE = 0 # Anomaly score : 0 for 1-MSP, 1 for H(u;.)
```

# Outlier exposure: code review

- Dataloader

```
train_dataloader = DataLoader(train_dataset, batch_size=BATCH, shuffle=True)
val_dataloader = DataLoader(val_dataset, batch_size=BATCH)
test_dataloader = DataLoader(test_dataset, batch_size=BATCH)
```

```
# outlier dataset
```

```
{ outlier_idx = [i for i,v in enumerate(mnist_train) if v[1] in OUTLIER_NUM]
  outlier_dataset = Subset(mnist_train,outlier_idx)
  outlier_dataloader = DataLoader(outlier_dataset, batch_size=BATCH, shuffle=True)
```



# Outlier exposure: code review

- Anomaly score

```
▶ loss_fn = nn.CrossEntropyLoss(label_smoothing=LBSMOOTH) # CCE
✓ def anomaly_score(logits, mode=0): # mode = 0 :
✓     if mode == 0:
        # Anomaly score = 1 - MSP
        softmaxprob = torch.softmax(logits, dim=1)
        MSP = torch.max(softmaxprob, dim=1).values
        return torch.tensor(1) - MSP
✓     if mode == 1:
        # Anomaly score = -H(u,.)
        uni = torch.ones(logits.shape[0], logits.shape[1]).float().to(device) / W
        torch.tensor(len(NORMAL_NUM)).float().to(device) # make uniform distribution
        score = -F.cross_entropy(logits, uni, reduction='none') # -CCE (negative since anomaly)
        #print(score)
        return score
optimizer = torch.optim.Adam(model.parameters(), lr=1e-3) # Adam as optimizer
```

# Outlier exposure: code review

- Trainer 1

```
# Here, outlier_dataloader is added. More GPU memory is required for Outlier exposure.
def train(dataloader, outlier_dataloader, model, loss_fn, optimizer):
    model.train()
    size = len(dataloader.dataset)
    losses = []
    in_losses = []
    oe_losses = []
    batch = 0
    for indist, oedist in zip(dataloader, outlier_dataloader):
        # unzip tuple
        in_image, in_idx = indist
        oe_image, oe_idx = oedist

        # in distribution dataset (normal data)
        logits, _ = model(in_image.to(device))
        p = F.one_hot(in_idx, ORGCLASS_NUM)[: , NORMAL_NUM]
        p = p.float().to(device)
        in_dist_loss = loss_fn(logits, p)
        in_losses.append(in_dist_loss.cpu().detach())

        # logit (batch, class)

        # each dist is tuple of (img, label)
```

make a tuple of two minibatches from train, outlier dataset

get (image, label) of training data

get (image, label) of outlier data

same as before

# Outlier exposure: code review

- Trainer 2

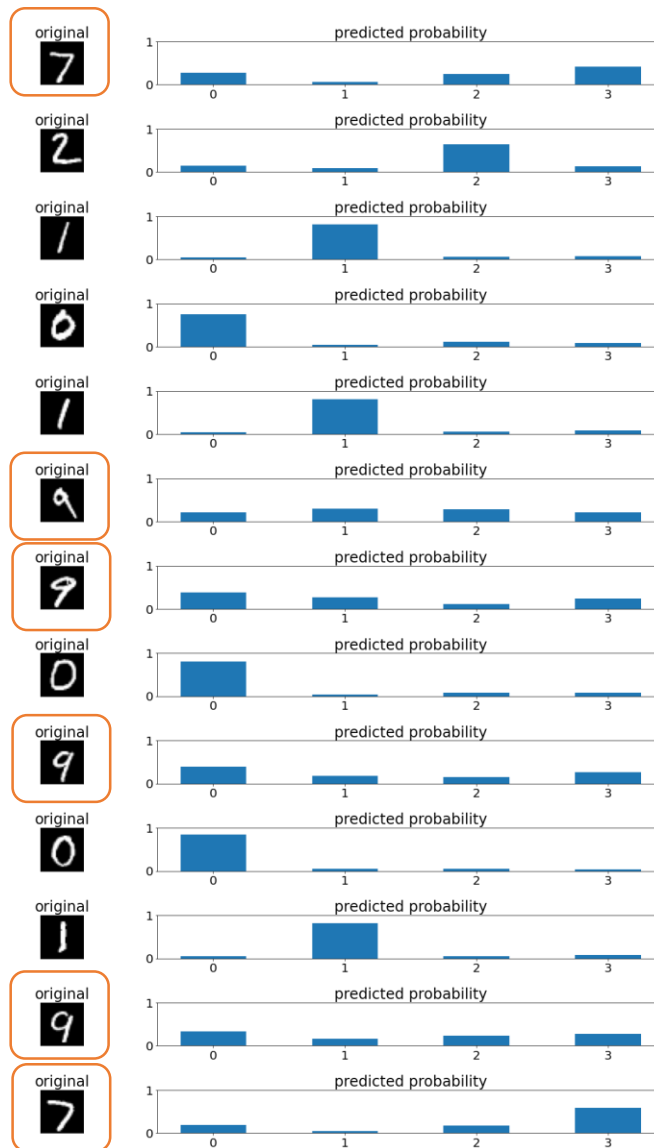


```
# OE outlier training
outlier_logits, _ = model(oe_image.to(device))
p_oe = torch.ones(oe_image.shape[0], len(NORMAL_NUM)) / torch.tensor(len(NORMAL_NUM)) # uniform dist.
p_oe = p_oe.float().to(device)
oe_loss = loss_fn(outlier_logits, p_oe) # CCE outlier loss
oe_losses.append(oe_loss.cpu().detach())

# joint training
    weighted sum of two losses
loss = in_dist_loss + 0.5*oe_loss # 0.5 for visual task
losses.append(loss.cpu().detach())

optimizer.zero_grad()
loss.backward()
optimizer.step()
```

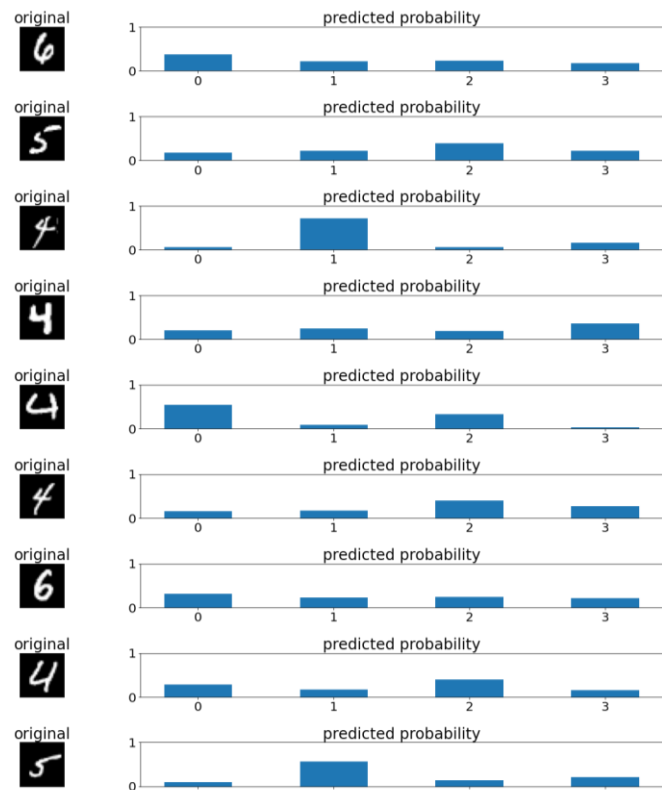
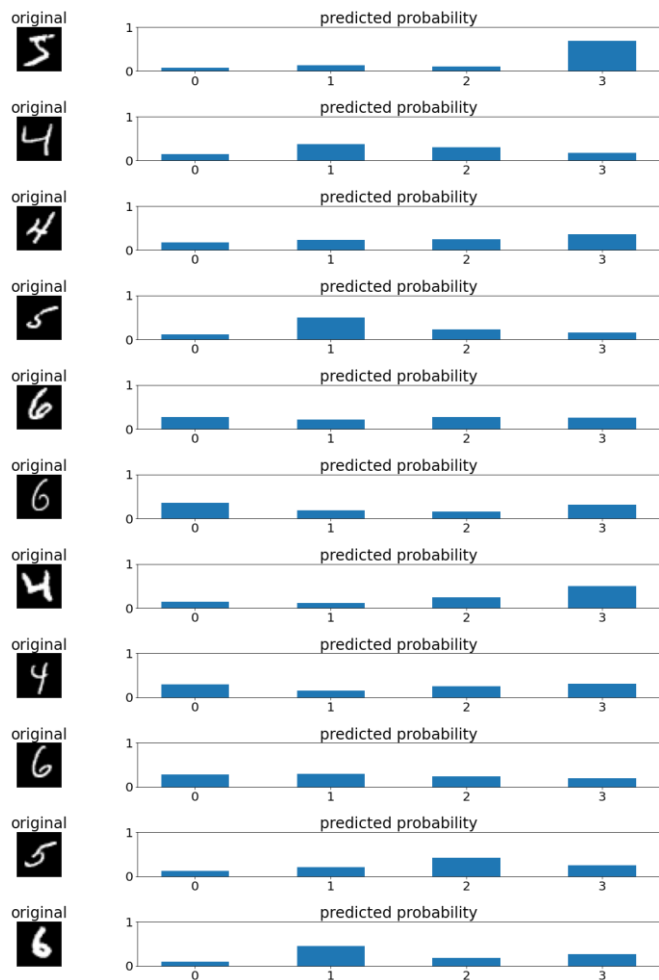
← target distribution: uniform

# Outlier exposure: result




# Outlier exposure: result

- Outlier prediction probability (seen outlier data during training)



# Outlier exposure: result

- t-SNE plot




# Other pretext tasks

1. Reconstruction
2. Classification
3. Mixup

**mixup: Beyond Empirical Risk Minimization**  
<https://arxiv.org/abs/1710.09412>

# Overfitting problem

- Empirical risk minimization (ERM)



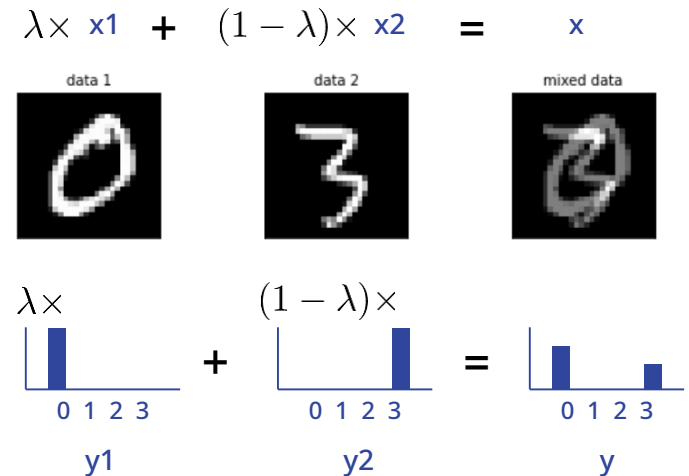
- DNN model minimizes expected risk at the observed points
  - Empirical risk minimization
- To minimize ERM, a model can memorize the training data
  - Vulnerable to unseen data
- What if we can find a **virtual feature-target pair** in the **vicinity** of the training pair?
  - Empirical vicinal risk minimization!



# Mixup

- Mix two classes
- Mix up both input and target probability

```
# y1, y2 should be one-hot vectors
for (x1, y1), (x2, y2) in zip(loader1, loader2):
    lam = numpy.random.beta(alpha, alpha)
    x = Variable(lam * x1 + (1. - lam) * x2)
    y = Variable(lam * y1 + (1. - lam) * y2)
    optimizer.zero_grad()
    loss(net(x), y).backward()
    optimizer.step()
```



- Not just an augmentation  
(target probability is redefined)
- Model is trained to predict mixed probability  $y$

# Mixup


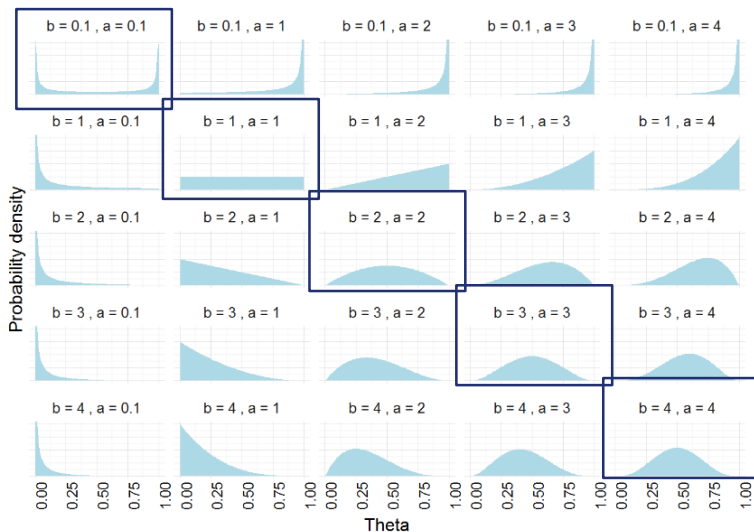
- Choice of  $\lambda$

- $\lambda$  is sampled from Beta distribution

$$\lambda \sim \text{Beta}(\alpha, \alpha)$$

- $\alpha$  controls the strength of interpolation

- Beta distribution (  $\text{Beta}(a,b)$  )



# Mixup: result

- MNIST dataset

- Without label mixing, Normal: [0,..., 3], Anomaly:[7, 8, 9]

